

# **Live Social Analyst**

The "Live AI" Engine for Real-Time Global Pulse

Pathway Datathon Submission

*Generated on: 2026-01-20 19:12*

# Live Social Analyst - Project Documentation

## 1. Documentation & Architecture

### High-Level Architecture:

Live Social Analyst is a real-time Retrieval-Augmented Generation (RAG) system powered by Pathway. It ingests live data from thousands of sources, processes it incrementally, and enables an LLM to answer questions about events happening right now.

### Core Components:

#### 1. Inputs:

- 34 OPML Files (1000+ RSS Feeds)
- NewsData.io API
- GNews API & HackerNews
- Twitter/X API (Ready)

#### 2. Pathway ETL Engine:

- Ingests streams in real-time
- Deduplicates identical stories
- Normalizes text and metadata
- Managing consistency and out-of-order data

#### 3. Storage & Retrieval:

- Hybrid Vector Store: In-memory live buffer for 0-latency access
- SQLite Archive: For historical persistence and complex filtering
- RAG Pipeline: Finds relevant context based on user queries

#### 4. AI & Inference:

- LLM: Gemini 1.5 Flash (Primary) / Groq Llama 3 (Fallback)
- Generates answers using the retrieved real-time context

### Deep Dive: The Life of a Data Point (Micro-Architecture):

1. T+0ms: News is published (RSS/API).
2. T+10ms: Pathway Connector (`pw.io`) detects the change and appends a row to the input Table.
3. T+15ms: The Streaming Engine propagates the update.
  - Deduplication Check: Global state checks if this URL hash was seen recently.
  - Normalization: HTML stripping and text cleaning (stateless transform).
4. T+100ms: Embedding: Using `pw.xpacks.llm`, the text is vectorized.
5. T+120ms: Indexing: The vector is upserted into the KNN index.
6. T+200ms: Ready for Query: The data is now live and retrievable by the RAG system.

## 2. Pathway Integration: Connectors & xPack

### 1. Live Data Ingestion with Pathway Connectors (pw.io):

- pw.io.fs.read: Used to monitor OPML files for real-time changes.
- pw.io.http.read: Used for connecting to NewsData.io and Twitter/X as infinite HTTP streams.
- pw.io.python.read: Used for wrapping custom Python generators as Pathway tables.

### 2. Streaming Transformations (pw.temporal & pw.state):

# Live Social Analyst - Project Documentation

- `pw.io.deduplicate(pathway.table, col=[url])`: Enforces uniqueness on the input stream.
- `pw.temporal.sliding(duration=5, step=1)`: Groups news items by 5-minute windows for "velocity" metrics.

## 3. LLM Integration (pw.xpacks.llm):

- `pw.xpacks.llm.embedders.GeminiEmbedder`: Real-time embedding of the text stream.
- `pw.xpacks.llm.retrievers.knn`: Live vector indexing and millisecond-latency retrieval for RAG.

## 3. Scalability, Observability & Extension

### Scalability:

- Rust-Based Engine: Pathway handles high-throughput streams (thousands of events/sec) efficiently on a single node.
- Distributed Ready: Although running locally for the demo, Pathway pipelines can scale to clusters using Docker/Kubernetes without code changes.

### Observability:

- Real-time logs track ingestion rates and LLM latency.
- "Global Pulse" UI visualizes the ingestion heartbeat.

### Extension to Other Domains:

- Finance: Monitor stock tickers and news for algorithmic trading signals.
- Supply Chain: Track shipping disruptions and weather events in real-time.
- Cybersecurity: Analyze log streams for anomaly detection.

## 4. Example Scenario: "The Market Crash"

### Scenario:

A major tech company announces a surprise acquisition at 10:00 AM.

1. T+0s (Announcement): The press release hits a news wire.
2. T+2s (Ingestion): Our OPML ingestor detects the new item.
3. T+3s (Processing): Pathway normalizes the text and updates the vector index.
4. T+5s (User Query): User asks "What just happened with [Company]?"
5. T+6s (Answer): The system retrieves the article ingested 3 seconds ago and generates: "Breaking: [Company] just announced an acquisition of..."

### Contrast with Traditional RAG:

A standard RAG system would wait for the nightly scrape (12-24 hours late), answering "I have no recent news on this."

## 5. Demonstration Pipeline

### The running demo illustrates:

1. "Global Pulse" Sidebar: Visualizing the raw stream of incoming news (updated every 5s via /pulse

# Live Social Analyst - Project Documentation

endpoint).

2. "Fetch Live" Button: Manually triggering a high-speed "Burst Mode" ingestion.
3. Chat Interface: Asking questions like "What is the latest on Apple?" and getting answers citing articles from 5 minutes ago.
4. Source Attribution: Every answer cites the specific RSS feed and timestamp, proving freshness.

## 6. Testing & Evaluation

Core validations performed:

1. Freshness Test: Verified that /pulse endpoint returns articles with timestamps < 5 minutes old.
2. Duplication Test: Confirmed that identical headlines from different sources are deduplicated.
3. Latency Test: Measured end-to-end time from "Ingest" to "Queryable" is < 2 seconds.
4. Load Test: Successfully handled ingestion of 1000+ feeds in "Burst Mode" without crashing.

## 7. Team Presentation Summary

Concept:

"Live Social Analyst" - A move from Static AI to Live AI.

Architecture:

A seamless pipeline: RSS/API -> Pathway Engine -> RAG -> Gemini.

Key Challenges:

- Ingesting noisy RSS feeds (solved with cleaning logic).
- Handling "Time": RSS feeds have inherent delays; we built a hybrid sorting mechanism to prioritize true publication time.
- Deduplication: Preventing the same AP/Reuters story appearing 50 times (solved with Pathway).

Why It Matters:

In a 24/7 world, static knowledge is a failure. This architecture proves that "Live AI" is accessible, scalable, and essential for decision-making.